

Appendix

Poster: *Do LLM judges of agent traces verify the process, or only the final outcome?*

Faizan · Trends in NLP · University of Trier · Examiner: Simon Munker

Repository: github.com/faizan10152/outcome-anchoring-llm-judges

A. References

Every entry, including the complete author list, was checked against the source record on 29 September 2026.

Advani, L. (2026). *From Confident Closing to Silent Failure: Characterizing False Success in LLM Agents*. arXiv:2606.09863.

Deshpande, D., Gangal, V., Mehta, H., Krishnan, J., Kannappan, A., & Qian, R. (2025). *TRAIL: Trace Reasoning and Agentic Issue Localization*. arXiv:2505.08638.

Lightman, H., Kosaraju, V., Burda, Y., Edwards, H., Baker, B., Lee, T., Leike, J., Schulman, J., Sutskever, I., & Cobbe, K. (2023). *Let's Verify Step by Step*. arXiv:2305.20050.

Liu, Y., Iter, D., Xu, Y., Wang, S., Xu, R., & Zhu, C. (2023). *G-Eval: NLG Evaluation using GPT-4 with Better Human Alignment*. Proceedings of EMNLP 2023. arXiv:2303.16634.

Lù, X. H., Kazemnejad, A., Meade, N., Patel, A., Shin, D., Zambrano, A., Stańczak, K., Shaw, P., Pal, C. J., & Reddy, S. (2025). *AgentRewardBench: Evaluating Automatic Evaluations of Web Agent Trajectories*. arXiv:2504.08942.

Mohammadi, H. (2026). *trajectory-judge: measuring what LLM judges miss when an agent reaches the right answer the wrong way*. Zenodo. <https://doi.org/10.5281/zenodo.21797927>

Yao, S., Shinn, N., Razavi, P., & Narasimhan, K. (2024). *τ -bench: A Benchmark for Tool-Agent-User Interaction in Real-World Domains*. arXiv:2406.12045.

Zhang, S., Yin, M., Zhang, J., Liu, J., Han, Z., Zhang, J., Li, B., Wang, C., Wang, H., Chen, Y., & Wu, Q. (2025). *Which Agent Causes Task Failures and When? On Automated Failure Attribution of LLM Multi-Agent Systems*. Proceedings of ICML 2025 (spotlight). arXiv:2505.00212.

Zheng, C., Zhang, Z., Zhang, B., Lin, R., Lu, K., Yu, B., Liu, D., Zhou, J., & Lin, J. (2024). *ProcessBench: Identifying Process Errors in Mathematical Reasoning*. arXiv:2412.06559.

Zheng, L., Chiang, W.-L., Sheng, Y., Zhuang, S., Wu, Z., Zhuang, Y., Lin, Z., Li, Z., Li, D., Xing, E. P., Zhang, H., Gonzalez, J. E., & Stoica, I. (2023). *Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena*. Advances in Neural Information Processing Systems 36 (NeurIPS 2023). arXiv:2306.05685.

Zhuge, M., Zhao, C., Ashley, D., Wang, W., Khizbullin, D., Xiong, Y., Liu, Z., Chang, E., Krishnamoorthi, R., Tian, Y., Shi, Y., Chandra, V., & Schmidhuber, J. (2024). *Agent-as-a-Judge: Evaluate Agents with Agents*. arXiv:2410.10934.

Software and tools

Ollama (2026). Local inference runtime, version as installed. Model weights: Qwen2.5 7B Instruct, Qwen3 14B and Llama 3.1 8B Instruct, obtained through the Ollama model library.

Python scientific stack: NumPy, pandas, SciPy, statsmodels and Matplotlib. Exact versions are pinned in `requirements.txt` in the repository.

B. Declaration on the use of generative AI tools

This section is provided in fulfilment of the requirement in the Declaration of Academic Integrity that generative AI tools be indicated with their product name and listed in full in an appendix.

Product name	Where and how it was used
Claude Code (Anthropic), model Claude Opus 5	Used as a programming and drafting assistant throughout the project: implementing the task environment, oracle, agent loop, fault injector, judge runner and analysis code; producing the figures; drafting and laying out the poster text; and locating and checking bibliographic records.

Scope and responsibility

- All numerical results reported on the poster were produced by executing the code in the repository against local language models. No result, table or figure value was written or estimated by a generative model.
- The experimental design, the research question and the three hypotheses were set by the author before implementation.
- Every bibliographic entry in Section A was verified by the author against the source record; none was accepted on the assistant's assertion alone.
- The author has read, checked and proofread the whole submission and takes full responsibility for its content, including any passages produced with assistance from the tool named above.

Reproducibility

The repository contains the complete environment, all generated agent traces, all 606 judge outputs, the analysis code and the figure scripts, so that every number on the poster can be recomputed from the committed data. The full development history, including the assistant's contributions, is visible in the commit log.

C. Declaration of Academic Integrity

The signed declaration follows on the next page.